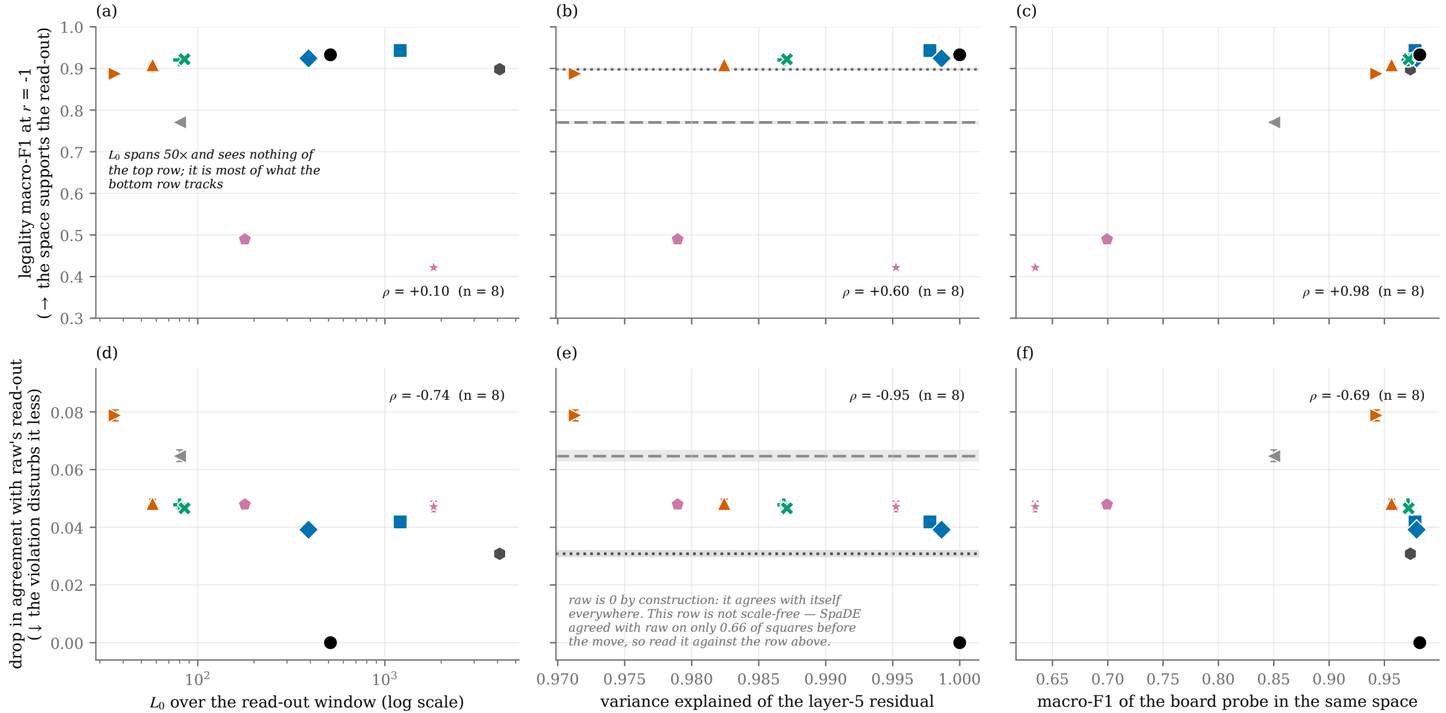*$L_0$ and variance explained are measured on this study's own read-out window (results/space_quality.csv); raw sits at $L_0$ = 512 and VE = 1 by construction, and the two untrained controls have no reconstruction objective, so on (b) and (e) each carries a horizontal level instead of an x position. The sparsity-matched null sits at exactly TopK's $L_0$ in (a) and (d). bars: 95% bootstrap CI over source games (n = 2000), and the same interval shades each control level. $\rho$ is Spearman over the eight trained dictionaries only and is descriptive at that n. The top row is a level, which this study cannot claim as a result; the panels are about which x tracks it.*