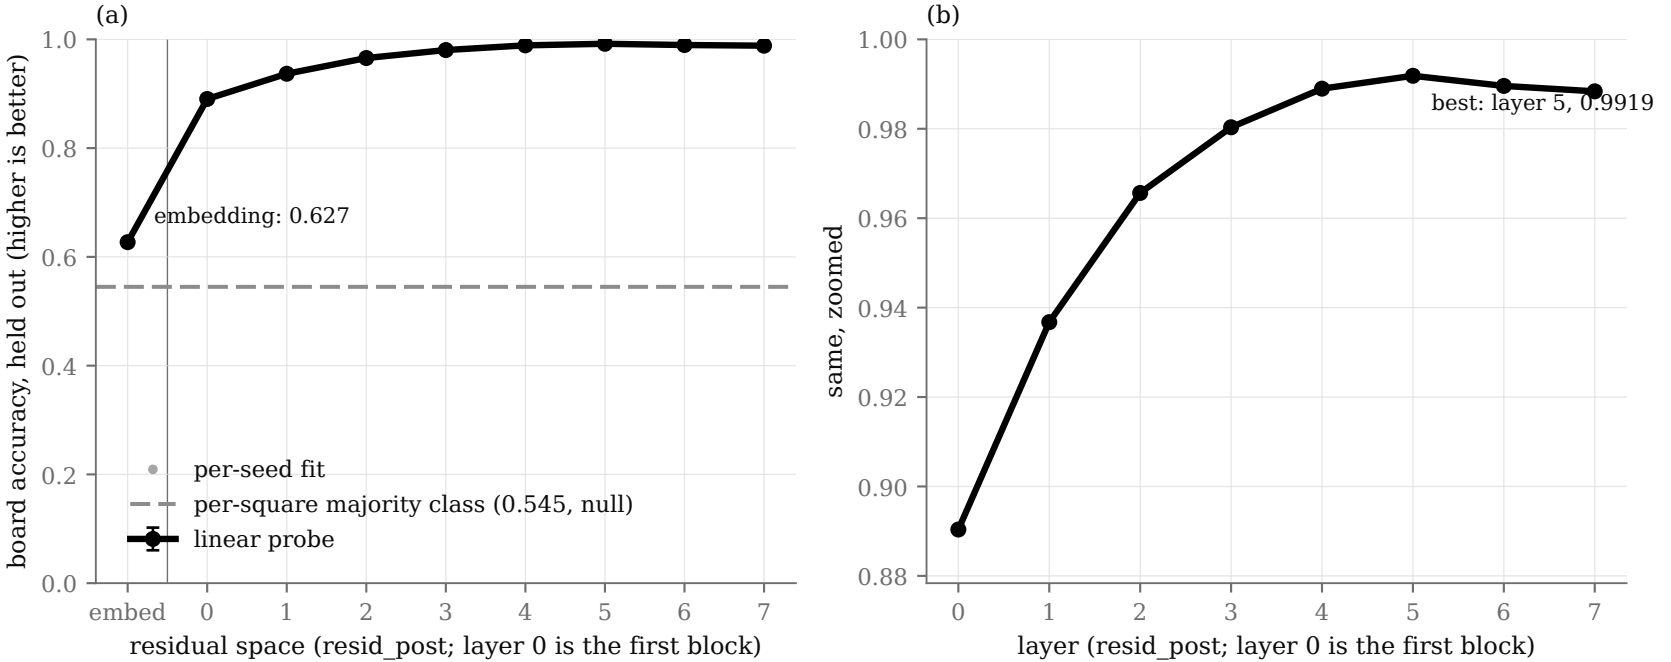

Bars: ±1 SD over 5 probe-initialisation seeds (largest 0.0001), drawn behind their mean. The games, the 8000/2000 split and the activations are fixed, so the bars are optimisation noise and not sampling noise over games. Accuracy is over 2000 held-out positions × 64 squares, 3 classes {empty, theirs, mine}. (b) drops the embedding space and truncates the y-axis; the marks carry intervals, so a truncated axis cannot exaggerate a difference the way a cut bar baseline does.